



# NVFP4 GEMM on NVIDIA B300


4,096 × 4,096 × 4,096




8,192 × 8,192 × 8,192



16,384 × 16,384 × 16,384



32,768 × 32,768 × 32,768



Throughput (PFLOP/s)

Throughput (PFLOP/s)

† BF16 output; other rows FP16. ThunderKittens pads K to a multiple of 768.  
CUPTI · CUDA graphs · cold L2